

Virtual Memory

Fan Wu

Department of Computer Science and Engineering


Shanghai Jiao Tong University

Spring 2022

Background

- Code needs to be in memory to execute, but entire program rarely used
 - Error code, unusual routines, large data structures
- Entire program code not needed at the same time
- Consider ability to execute partially-loaded program
 - Program no longer constrained by limits of physical memory
 - Program could be larger than physical memory

Virtual Memory That is Larger Than Physical Memory




Virtual Memory

- **Virtual Memory** – separation of user logical memory from physical memory
 - Only part of the program needs to be in memory for execution
 - Logical address space can therefore be much larger than physical address space
 - Allows memory address spaces to be shared by several processes
 - Allows for more efficient process creation
 - More programs running concurrently
 - Less I/O needed to load or swap processes
- Virtual memory can be implemented via:
 - Demand paging
 - Demand segmentation

Demand Paging

- Could bring entire process into memory at load time
- Or bring a page into memory only when it is needed
 - Less I/O needed, no unnecessary I/O
 - Less memory needed
 - Faster response
 - More users
- Page is needed \Rightarrow reference to it
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory
- **Lazy swapper (pager)** – never swaps a page into memory unless page will be needed

Swap Paged Memory to Disk Space



Valid-Invalid Bit

- With each page table entry a **valid-invalid** bit is associated (**v** \Rightarrow in-memory – **memory resident**, **i** \Rightarrow not-in-memory)
- Initially, valid-invalid bit is set to **i** on all entries
- Example of a page table snapshot:


Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table

page fault \downarrow
页面不在内存里
要从磁盘调入。

- During address translation, if valid-invalid bit in page table entry is **i** \Rightarrow page fault

Page Table with Pages Not in Main Memory



Page Fault

- If there is a reference to a page and the page is not in memory, the reference will trap to operating system:

page fault

1. Operating system looks at page table to decide:

- Invalid reference \Rightarrow abort

- Just not in memory \rightarrow 如果是不在内存, 那么就创建一个新的页, 然后把这个页所有在 page table 的对应位置中

2. Get empty frame

3. Swap page into frame via scheduled disk operation


4. Reset tables to indicate page now in memory

Set validation bit = **v**

\rightarrow 之后再重新执行这个取址操作

5. Restart the instruction that caused the page fault


Steps in Handling a Page Fault



What Happens if There is no Free Frame?

- Page replacement – find some page in memory, but not really in use, page it out
 - Algorithm – terminate? swap out? replace the page?
 - Performance – want an algorithm which will result in minimum number of page faults

Page Replacement



Page Replacement

- Prevent over-allocation of memory by modifying page-fault service routine to include page replacement
- Use **modify (dirty) bit** to reduce overhead of page transfers – only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory

page 0 将被替换 } $\text{modify bit} = 0$ 说明与 disk 一样, 直接换掉
 $\text{modify bit} = 1$ 与 disk 不同, 要写回 disk, 再换上 new page.

Page Replacement Algorithms

■ Page-replacement algorithm

- Want lowest page-fault rate on both first access and re-access

■ Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string

- String is just page numbers, not full addresses
- Repeated access to the same page, which is still in memory, does not cause a page fault

■ In all our examples, the reference string is

7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1

Page-Replacement Algorithms


- First-In-First-Out (FIFO) Page Replacement
- Optimal Page Replacement 最优页面替换
- Least Recently Used (LRU) Page Replacement
- LRU Approximation Page Replacement 二次机会.
- Counting Page Replacement $\begin{cases} \text{MFU} \\ \text{LFU} \end{cases}$

FIFO Page Replacement

- When a page must be replaced, the oldest page is chosen.

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

(不会因为重复访问而刷新)


- Page faults: 15
- Consider the following reference string:
0 1 2 3 0 1 2 3 0 1 2 3

Optimal Page Replacement

- Replace page that will not be used for longest period of time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

- Page faults: 9
- How do you know this?
 - Can't read the future
- Used for measuring how well your algorithm performs

Least Recently Used (LRU) Page Replacement

最近最久未使用算法

- Use past knowledge rather than future
- Replace page that has not been used in most amount of time
- Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used

LRU Approximation Algorithms

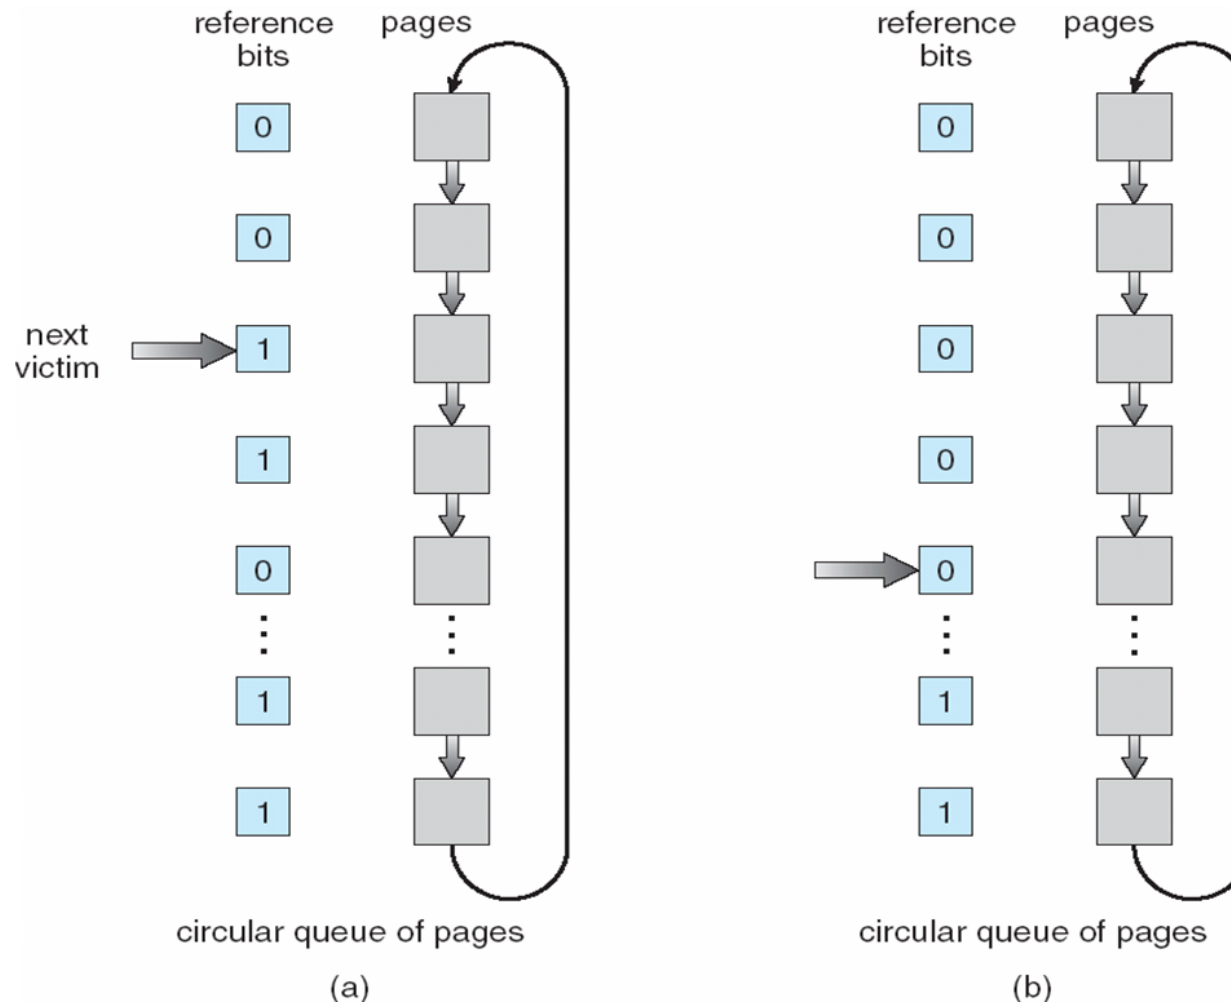
■ Reference bit/ byte

- With each page associate a bit, initially = 0
- When page is referenced, bit set to 1
- Replace any with reference bit = 0 (if one exists)
 - ▶ We do not specify the order, however

■ Second-chance algorithm

- Generally FIFO, plus hardware-provided reference bit
- Circular replacement
- If page to be replaced has
 - ▶ Reference bit = 0 -> replace it
 - ▶ Reference bit = 1 then:
 - set reference bit 0, leave page in memory
 - replace next page, subject to same rules

Second-Chance Algorithm



- three** frames.

0 1 2 3 2 3 0 4 5 2 3 1 4 3 2 6 3 2 1 2

Draw a diagram to show the page replacement using **Second-Chance Algorithm** and calculate the number of page faults.

0 1 2 3 2 3 0 4 5 2 3



Counting Algorithms

- Keep a counter of the number of references that have been made to each page
- **Least Frequently Used (LFU) Algorithm**: replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used
- Not commonly used

Homework

- Reading
 - Chapter 9
- Exercise
 - See course website

Demand Paging

System Characteristics	
Size of memory	16 bytes
Frame Size	4 bytes per frame <i>7 4 frames</i>
Memory Management Structure	Inverted Page Table
Replacement Policy	LRU, Global Replacement
Virtual Page Size	4 bytes per page
Logical Addressing Space Size	32 bytes <i>= 2⁵</i> <i>7 8 pages</i>
Backing Store Size	12 blocks
Backing Store Block Size	4 bytes per block

Process Table

Process ID	0	1	2
Process Size (Bytes)	12	14	13
Pages allocated	3	4	4
Backing Store Map (Page → Block)			
Page 0	BS 0	BS 3	BS 7
Page 1	BS 1	BS 4	BS 8
Page 2	BS 2	BS 5	BS 9
Page 3		BS 6	BS 10

System Snapshot

Main Memory

Address	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Contents	O	T	O	N	G	F	U	N	-	-	-	-	A	D	*	F

Backing Store

Block	0	1	2	3	4	5	6	7	8	9	10
Contents	THRE	AD*F	UN--	RATE	*MON	OTON	IC--	DEMA	ND*P	AGIN	G---

Inverted Page Table

Frame	Page #	PID	Valid Bit	Ref Word (Low = older)	Modified Bit
0	2	1	T	2	F
1	3	2	T	1	T
2	-	-	F	-	-
3	1	0	T	3	F

PID 0 : Write 'A' at logical memory Address 11

Process ID	0
Process Size (Bytes)	12
Pages allocated	3
Backing Store Map (Page → Block)	
Page 0	BS 0
Page 1	BS 1
Page 2	BS 2
Page 3	

Main Memory

Address	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Contents	O	T	O	N	G	F	U	N	-	-	-	-	A	D	*	F
Change									U	N	-	A				

Backing Store

Block	0	1	2	3	4	5	6	7	8	9	10
Contents	THRE	AD*F	UN--	RATE	*MON	OTON	IC--	DEMA	ND*P	AGIN	G---
Change											

Inverted Page Table

Frame	VP #	PID	Valid Bit	Ref Word (Low = older)	Modified Bit
0	2	1	T	2	F
1	3	2	T	1	T
2	-	2	F	T	-
3	1	0	T	3	F

PID 1 : Read logical memory Address 6

Process ID	1
Process Size (Bytes)	14
Pages allocated	4
Backing Store Map (Page → Block)	
Page 0	BS 3
Page 1	BS 4
Page 2	BS 5
Page 3	BS 6

Main Memory

Address	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Contents	O	T	O	N	G	F	U	N	-	-	-	-	A	D	*	F
Change					*	M	O	N	U	N	-	A				

Backing Store

Block	0	1	2	3	4	5	6	7	8	9	10
Contents	THRE	AD*F	UN--	RATE	*MON	OTON	IC--	DEMA	ND*P	AGIN	G---
Change											GFUN

Inverted Page Table

Frame	VP #	PID	Valid Bit	Ref Word (Low = older)	Modified Bit
0	2	1	T	2	F
1	3	1	T	1	T
2	-	2	F	-	-
3	1	0	T	3	F